

Machine Learning Engineer – Mid-Level

Seattle, WA | Hybrid (2-3 days onsite) | Full-time | Department: Personalization

Reports to	Staff ML Engineer	Team	Recommendations & Ranking
Level	Mid-Level	Travel	Up to 10%
Posting date	February 24, 2026	Start window	Next 6-10 weeks

About the company

Aster Retail Labs builds e-commerce personalization software. Our teams ship software frequently, measure outcomes, and invest in reliability and security. We build ranking and recommendation services that power product discovery across web and mobile experiences.

Role summary

Own ML components end-to-end: problem framing, model development, experimentation, and reliable production deployment with monitoring and iteration.

What you will do

- Define training datasets, baselines, metrics, and offline/online evaluation plans
- Develop models and pipelines for batch and real-time inference; optimize latency and cost
- Partner with product and data stakeholders to ship measurable improvements (A/B tests where applicable)
- Build or improve MLOps workflows: CI/CD for models, automated retraining triggers, monitoring and alerting
- Review peers' code and mentor junior engineers on best practices
- Contribute to documentation, runbooks, and incident response for ML services

What we are looking for

Required

- 3+ years of experience building ML systems or equivalent project track record
- Strong Python and at least one deep learning framework (PyTorch or TensorFlow)
- Hands-on experience with production deployment (APIs, batch jobs, streaming, or edge)
- Working knowledge of data pipelines and storage (SQL, Spark, warehouses/lakes)
- Ability to debug model and data issues using logs, metrics, and experiments

Preferred

- Experience with feature stores, model registries, and monitoring tools

- Experience with distributed training or GPU optimization
- Familiarity with responsible AI practices (privacy, bias testing, interpretability)

Tech stack (typical)

Python, PyTorch/TensorFlow, Spark, SQL, Docker/Kubernetes, MLflow, Airflow, AWS/GCP; Feature store, Redis, A/B testing platform

Success looks like (first 90 days)

- Deliver an end-to-end model improvement with an agreed evaluation plan and measurable lift
- Harden a training or inference pipeline (tests, monitoring, runbooks) and reduce operational toil
- Partner with product to define a rollout plan; participate in an A/B test or controlled launch
- Mentor a junior engineer through at least one project milestone

Compensation & benefits

Base salary range: \$145,000 - \$185,000 per year. Includes eligibility for bonus depending on level. Final offer depends on experience, skills, and location.

- Competitive salary and annual bonus or equity (role dependent)
- Medical, dental, and vision coverage
- 401(k) or local retirement plan with employer match (where applicable)
- Flexible PTO and paid holidays
- Learning budget for courses, conferences, and books
- Home-office/remote-work stipend (where applicable)

Interview process

- Intro call with Talent Partner (20-30 min)
- Technical screen (coding + ML/system fundamentals)
- Onsite/virtual loop (2-4 interviews: ML, systems, product, collaboration)
- Final conversation with hiring manager

We are an equal opportunity employer. All qualified applicants will receive consideration without regard to race, color, religion, sex, sexual orientation, gender identity, national origin, disability, veteran status, or any other legally protected status.

If you require an accommodation during the application or interview process, please let us know. We will work with you to meet your needs.