

Machine Learning Engineer – Senior

San Francisco, CA | Hybrid (3 days onsite) | Full-time | Department: Threat Detection

Reports to	Director of Engineering	Team	Detection ML
Level	Senior	Travel	Up to 15%
Posting date	February 24, 2026	Start window	Next 6-12 weeks

About the company

CloudForge Security builds cybersecurity products for anomaly detection. Our teams ship software frequently, measure outcomes, and invest in reliability and security. We detect suspicious behavior across large volumes of logs and events, and we ship models with strict uptime and low false-positive requirements.

Role summary

Lead technical delivery of ML projects with significant business impact. You will design architectures, set standards for quality and reliability, and mentor engineers across teams.

What you will do

- Architect scalable ML systems (training, feature pipelines, inference, monitoring) for multiple products
- Set evaluation standards and guardrails; ensure reproducibility, reliability, and safe rollout strategies
- Drive cross-functional execution: align product, data, and engineering on milestones and success metrics
- Optimize models and infrastructure for performance, latency, and cost; establish capacity plans
- Mentor engineers; lead design reviews and post-incident analyses with action-oriented follow-ups
- Influence roadmap: identify high-leverage opportunities and de-risk technical choices

What we are looking for

Required

- 6+ years of experience in ML engineering, data-driven product engineering, or similar
- Proven track record shipping ML models to production and maintaining them over time
- Strong systems thinking: APIs, distributed systems, observability, and failure modes
- Deep understanding of ML evaluation and experimentation (offline metrics, A/B tests, monitoring drift)
- Excellent communication and ability to lead without authority

Preferred

- Experience operating ML services at scale (high QPS, multiple regions, strict SLAs)

- Experience with retrieval, ranking, personalization, or forecasting systems
- Experience designing governance for model risk, privacy, and compliance

Tech stack (typical)

Python, PyTorch, Kubernetes, Kafka, Spark, Feature Store, Model Registry, Observability stack; Kafka, ClickHouse, Kubernetes, SLOs

Success looks like (first 90 days)

- Lead design and delivery of a high-impact AI/ML initiative with clear success metrics
- Improve reliability of an ML service (alerts, dashboards, SLOs) and reduce incident rate
- Establish a repeatable evaluation framework used by multiple engineers/teams
- Mentor engineers through reviews and design guidance; raise the quality bar

Compensation & benefits

Base salary range: \$190,000 - \$245,000 per year. Includes eligibility for bonus/equity depending on level. Final offer depends on experience, skills, and location.

- Competitive salary and annual bonus or equity (role dependent)
- Medical, dental, and vision coverage
- 401(k) or local retirement plan with employer match (where applicable)
- Flexible PTO and paid holidays
- Learning budget for courses, conferences, and books
- Home-office/remote-work stipend (where applicable)

Interview process

- Intro call with Talent Partner (20-30 min)
- Technical screen (coding + ML/system fundamentals)
- Onsite/virtual loop (2-4 interviews: ML, systems, product, collaboration)
- Final conversation with hiring manager

We are an equal opportunity employer. All qualified applicants will receive consideration without regard to race, color, religion, sex, sexual orientation, gender identity, national origin, disability, veteran status, or any other legally protected status.

If you require an accommodation during the application or interview process, please let us know. We will work with you to meet your needs.