

Senior AI Engineer

Toronto, ON | Hybrid (2-3 days onsite) | Full-time | Department: Applied AI

Reports to	Director of AI	Team	Operations Copilot
Level	Senior	Travel	Up to 15%
Posting date	February 24, 2026	Start window	Next 6-12 weeks

About the company

FoundryOps builds industrial operations software for manufacturing teams. Our teams ship software frequently, measure outcomes, and invest in reliability and security. We build copilots that help engineers troubleshoot equipment, summarize logs, and recommend actions with grounded retrieval and robust evaluation.

Role summary

Own high-impact AI initiatives and set technical direction for how we build, evaluate, and operate AI systems. You will lead designs, mentor others, and drive reliable delivery across teams.

What you will do

- Lead architecture and implementation of AI systems (LLMs, retrieval, model serving, data pipelines)
- Define evaluation frameworks and quality gates; create repeatable test suites and dashboards
- Build guardrails for safety, privacy, and compliance; partner with security and legal stakeholders
- Drive performance and cost optimization across inference, retrieval, and orchestration
- Mentor engineers and review designs; establish patterns and standards for the AI platform
- Own incident response and reliability plans for AI services with clear SLAs

What we are looking for

Required

- 6+ years of engineering experience with 3+ years focused on ML/AI systems
- Proven track record delivering production AI features with measurable outcomes
- Strong backend engineering and systems design skills
- Experience with evaluation, experimentation, and monitoring for AI quality
- Strong communication and stakeholder management

Preferred

- Experience with fine-tuning, distillation, or optimizing LLM inference

- Experience leading technical roadmaps and mentoring across multiple teams
- Experience with regulated domains (finance, healthcare) or strict privacy constraints

Tech stack (typical)

Python, PyTorch, LLM tooling, Retrieval/Vector DB, Kubernetes, Data pipelines, Observability; Orchestration framework, vector search

Success looks like (first 90 days)

- Lead design and delivery of a high-impact AI/ML initiative with clear success metrics
- Improve reliability of an ML service (alerts, dashboards, SLOs) and reduce incident rate
- Establish a repeatable evaluation framework used by multiple engineers/teams
- Mentor engineers through reviews and design guidance; raise the quality bar

Compensation & benefits

Base salary range: \$185,000 - \$240,000 per year. Includes eligibility for bonus/equity depending on level. Final offer depends on experience, skills, and location.

- Competitive salary and annual bonus or equity (role dependent)
- Medical, dental, and vision coverage
- 401(k) or local retirement plan with employer match (where applicable)
- Flexible PTO and paid holidays
- Learning budget for courses, conferences, and books
- Home-office/remote-work stipend (where applicable)

Interview process

- Intro call with Talent Partner (20-30 min)
- Technical screen (coding + ML/system fundamentals)
- Onsite/virtual loop (2-4 interviews: ML, systems, product, collaboration)
- Final conversation with hiring manager

We are an equal opportunity employer. All qualified applicants will receive consideration without regard to race, color, religion, sex, sexual orientation, gender identity, national origin, disability, veteran status, or any other legally protected status.

If you require an accommodation during the application or interview process, please let us know. We will work with you to meet your needs.