

Principal AI/Machine Learning Engineer

New York, NY | Hybrid (2-3 days onsite) | Full-time | Department: AI Platform

| | | | |
|--------------|-------------------|--------------|--------------------------|
| Reports to | CTO | Team | AI Platform Architecture |
| Level | Principal | Travel | Up to 25% |
| Posting date | February 24, 2026 | Start window | Next 10-16 weeks |

About the company

Atlas Commerce builds commerce platform and AI infrastructure for global sellers. Our teams ship software frequently, measure outcomes, and invest in reliability and security. You will set standards for how we build and operate ML across personalization, search, risk, and generative AI products.

Role summary

Drive architecture and technical strategy for company-wide AI and ML systems. You will influence multiple teams, set engineering standards, and tackle the highest-impact technical problems.

What you will do

- Define long-term architecture for ML/AI platforms, including data, training, serving, and evaluation
- Lead high-stakes projects end-to-end; de-risk decisions through prototypes and technical reviews
- Set standards for reliability, security, privacy, and responsible AI across the organization
- Mentor senior engineers and leads; grow technical leadership capacity across teams
- Partner with executives on roadmap, resourcing, and build-vs-buy decisions
- Represent the organization in technical forums; contribute to hiring and calibration

What we are looking for

Required

- 10+ years of engineering experience with deep expertise in ML/AI systems
- History of influencing architecture across multiple teams and products
- Strong software engineering fundamentals and distributed systems knowledge
- Track record of building scalable, maintainable platforms and setting standards
- Excellent communication and ability to drive alignment across stakeholders

Preferred

- Experience with large-scale ranking/retrieval, multi-modal models, or personalization platforms
- Experience with privacy-preserving ML or federated learning

- Experience with cost management for large inference workloads

Tech stack (typical)

Python, ML platform stack, Distributed systems, Kubernetes, Data lake/warehouse, Observability; Multi-tenant platform, governance, cost controls

Success looks like (first 90 days)

- Define and socialize a durable architecture for AI/ML platform components across teams
- Unblock a critical initiative by de-risking the toughest technical decision(s)
- Set organization-wide standards for evaluation, reliability, and responsible AI
- Mentor technical leads and contribute to hiring and calibration

Compensation & benefits

Base salary range: \$250,000 - \$330,000 per year. Includes eligibility for bonus/equity depending on level. Final offer depends on experience, skills, and location.

- Competitive salary and annual bonus or equity (role dependent)
- Medical, dental, and vision coverage
- 401(k) or local retirement plan with employer match (where applicable)
- Flexible PTO and paid holidays
- Learning budget for courses, conferences, and books
- Home-office/remote-work stipend (where applicable)
- Executive wellness benefit and expanded equity program

Interview process

- Intro call with Talent Partner (20-30 min)
- Technical screen (coding + ML/system fundamentals)
- Onsite/virtual loop (2-4 interviews: ML, systems, product, collaboration)
- Final conversation with hiring manager
- Executive alignment conversation (30 min)

We are an equal opportunity employer. All qualified applicants will receive consideration without regard to race, color, religion, sex, sexual orientation, gender identity, national origin, disability, veteran status, or any other legally protected status.

If you require an accommodation during the application or interview process, please let us know. We will work with you to meet your needs.