

About the job

First ML Engineer hire. Build the Data and ML flywheel.

The problem

Compliance in MedTech is complex, manual, and error-prone, but it's also a differentiator, a moat, and a strategic endeavor. It's like wind in sailing: work with it and it propels you forward; ignore it and you're fighting the current.

Klaris is an AI platform transforming how MedTech companies navigate regulatory documentation and approvals. We help medical device companies gain faster approvals, pass audits, and focus on patients and innovation while giving them confidence in their decisions.

About us

We're a pre-seed AI startup in London, backed by Meridian Health Ventures, Antler, Vento Ventures, and incredible angels with a runway to scale. Francesco (CEO & co-founder) spent a decade in MedTech bringing multiple lifesaving products from concept to patients, he's lived the compliance pain we're solving. Mihai (CTO & co-founder) has a PhD on decision making in multi-agent systems and led research in safe and interactive planning for self-driving.

The role

As our Founding ML Engineer, you'll work on our Data and ML engine necessary to develop custom, state-of-the-art algorithms to automate safety-critical regulatory tasks in Life Sciences. This means developing robust, scalable, streamlined infrastructure and frameworks covering all aspects of DataOps and MLOps.

You'll take ownership of critical components, and work closely with an experienced research scientist to establish the foundations that we will scale on. As a Founding member of our team, you will not just impact the technical roadmap, but also the team's culture, and technical hiring strategy.

What we're looking for

You should have a deep ML expertise in developing frameworks and algorithms as well as turning them into successful products. You should be able to quickly prototype to test hypotheses, and also convert the latest developments into robust and scalable products. We don't care about degrees or years on a CV; we want to see evidence of exceptional ML engineering craft.

Startup DNA is essential: comfort with ambiguity, high agency, and excitement about 0-to-1 building. We're looking for someone who thrives with autonomy and makes independent decisions without needing structure or oversight.

While experience in the MedTech or other Regulated industry is a strong bonus, this should not put you off from applying. We foremost value fast learners, determined and adaptable candidates.

Core technical skills

Strong competence and proven experience (4+ years) with:

- **MLOps**: automating the ML lifecycle, including CI/CD and monitoring, with familiarity with tools (MLflow, Weights & Biases or similar).
- **Data Pipeline**: building and managing data pipelines (e.g. ETL/DAGs), including document processing tasks and/or Optical Character Recognition (OCR).
- **Natural Language Processing**: retrieval, knowledge graphs, tokenization, chunking, named entity recognition etc.
- **LLMs/SLMs**: productionising LLMs/SLMs into agentic, RAG, multi-model or multi-agent solutions.
- **SW Languages**: Python, and a basic understanding of Javascript/TypeScript.

Hands-on experience:

- **Infrastructure**: cloud platforms (GCP or AWS), containerization tools (Kubernetes, Docker), and infrastructure-as-code (Terraform).
- **Databases**: Postgres, Neo4j and/or vector databases.
- **Backend Development Frameworks**: FastAPI, Pydantic, and SQLAlchemy.
- **LLM Development Frameworks**: e.g. DSPy, LangChain, etc.

What we offer

- Compensation: competitive base salary + equity, with flexibility to adjust the cash/equity mix for more senior candidates.
- Time off: 25 days + UK bank holidays.
- Environment: Office in London (hybrid). We value in-person collaboration and would love you in our Aldgate East office regularly, especially initially. But we're also parent-friendly and offer genuine flexibility.
- Benefits: Health insurance and pension schemes
- Visa sponsorship: We can sponsor visas for the right candidate.

Target start: As soon as possible

Interview Process

- 30-min chat with CTO
- 1h-1.5h technical deep-dive
- 45-min lunch or coffee with founders
- Reference Checks + Offer