

# About the job

Lead AI Engineer

Remote-first (UK - Ideally Scotland) with occasional travel to client sites and team meet-ups

£100,000-£120,000 + Benefits

**Travel:** Approx. 1 day per week on average

MBN have partnered exclusively with an established start up Data Consultancy.

We're looking for an experienced, hands-on **Lead AI Engineer** to design, build, and scale **production-grade, AI-native systems** used by a global group of companies.

This is a senior technical role for someone who combines strong software engineering fundamentals with deep expertise in modern AI frameworks and LLM-driven architectures.

You'll shape AI architecture strategy, lead delivery of real-world systems, and mentor engineers – all while staying close to the code.

This role is about **building real systems in production**, not just PoCs.

## What You'll Be Doing

- Design and build AI-native applications and intelligent workflows in Python
- Architect and deliver LLM-powered systems, agents, RAG pipelines, and orchestration layers
- Lead production deployments using modern AI tooling, including:
- LangChain, FastMCP, LlamaIndex
- OpenAI / Anthropic / Gemini SDKs
- Vector databases (Pinecone, Weaviate, Chroma, FAISS)
- Design multi-agent systems and tool-augmented workflows
- Establish best practices for:
- Prompt engineering
- Evaluation and testing of AI systems
- Observability, monitoring, and performance optimisation
- Provide technical leadership and architectural oversight
- Mentor engineers and raise the bar on AI engineering practices
- Work closely with stakeholders to turn real business problems into scalable AI solutions

## Must-Have

- 7+ years of software engineering experience (strong Python)
- 3+ years building AI/ML systems in production
- Hands-on experience with:
- LangChain and/or LlamaIndex

- FastMCP or similar model context protocol tools
- OpenAI, Anthropic, Gemini, or similar APIs
- Strong understanding of:
  - LLM capabilities and limitations
  - RAG, embeddings, vector search
  - Tool calling, function calling, and guardrails
  - Secure, production-grade AI system design
  - Experience integrating vector databases into real systems
  - Solid engineering fundamentals (CI/CD, testing, design patterns, code review)
  - Experience with Docker and cloud platforms (preferably Azure)

#### Nice to Have

- Multi-agent system design and deployment
- LLM evaluation frameworks (LangSmith, DeepEval, custom eval pipelines)
- ML lifecycle tooling (MLflow, Weights & Biases)
- Exposure to fine-tuning, LoRA, or model adaptation
- Experience replacing LLM APIs with offline SLMs at scale
- Startup or high-growth environment experience
- Driving licence and own transport#

#### Why Join?

- Work on **AI systems that actually ship and get used**
- High autonomy and genuine technical ownership
- Remote-first, flexible working
- Projects spanning modernisation, AI transformation, and AI-native platforms
- Founders and leaders who are deeply technical
- No dead-end PoCs – your work goes live

#### What's On Offer

- 10% discretionary annual bonus
- 35 days annual leave (including public holidays)
- Pension with 5% employer contribution
- Private healthcare scheme