

# About the job

## About Native

Native is a VC backed AI startup building the world's most accurate AI for databases reasoning. This is **hard**.

This is an unsolved problem that once solved will change how every enterprise operates.

We recently published our results, achieving rank 1 on the industry standard benchmark.

Our research team is based in London.

Everyone has high agency, obsessed with the problem we are solving, embraces the intensity of a 6-day week and has the ambition to build something world changing.

We are exceptional research engineers and scientists taking a novel approach to accurate database reasoning.

## About the role

This is a founding LLM agents research engineer role, with the opportunity to define how large language models act as reliable, long-horizon agents over complex enterprise data.

- Design and build LLM-powered agent architectures for reasoning over complex tabular, relational, and heterogeneous data
- Develop planning, memory, tool-use, and verification mechanisms that allow LLM agents to operate accurately over databases
- Research and implement techniques for grounding LLM agents in structured data representations (tables, graphs, schemas)
- Build training, simulation, and evaluation environments for long-horizon, multi-step agent behavior
- Develop self-supervised and reinforcement learning objectives for improving agent reliability and correctness
- Integrate foundation model embeddings and symbolic components into agent workflows
- Own agent systems end-to-end, from research prototypes to production-grade systems

## Required experience

- Proven ability to lead or significantly contribute to teams tackling hard research-engineering problems
- MSc / PhD in Mathematics, Physics, Computer Science, or a related field
- Strong hands-on experience building LLM-based systems with modern ML tooling (PyTorch, Ray, JAX)
- Experience designing agentic systems (planning, memory, tool use, multi-step reasoning)

- Strong understanding of transformers and foundation model behavior
- Comfortable owning whatever needs doing, from research ideas to infrastructure
- Experience optimizing large-scale training, fine-tuning, or simulation workloads on GPU clusters

## Why join us

This is a rare chance to join a founding team operating at the frontier of AI, with the resources to win.

- Impact: You will push the frontier of foundation models for an under-explored, massively valuable domain.
- Team: We are a small, collaborative, world-class research team based in central London.
- Founding level equity + competitive salary: We want you to share in the upside of what we are building. As part of the founding team, you will receive meaningful ownership alongside a salary that allows you to focus on building something great.
- Direction: Chance to shape the direction of a new class of foundation model from day one.

## Culture

- At Native, we look for people with exceptional research and engineering ability, who want to win.
- Everyone loves building and is ambitious enough to want to change how the world operates for the better.
- We run towards hard problems.
- We are creative problem solvers, who keep going until we find the underlying principle of the problem, and build to solve that.
- If you have an entrepreneurial spirit, are intrinsically motivated and want to have a huge impact, we are a great fit.